

Lanyue Zhang (606086099)  
Juntao Li (906092273)  
Wenxuan Du (506084251)  
Steven Li (705494229)

## Final Project

June 9, 2023

```
[1]: import pandas as pd
import numpy as np
import itertools

#Plotting libraries
import matplotlib.pyplot as plt
import seaborn as sns
import altair as alt

#statistics libraries
import statsmodels.api as sm
import statsmodels.formula.api as smf
import scipy
from scipy.stats import anderson
from statsmodels.tools.eval_measures import rmse
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import month_plot, seasonal_plot, plot_acf, \
    plot_pacf, quarter_plot
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.holtwinters import ExponentialSmoothing, SimpleExpSmoothing
from statsmodels.stats.diagnostic import acorr_ljungbox as ljung
from statsmodels.tsa.statespace.tools import diff as diff
import pmdarima as pm
from pmdarima import ARIMA, auto_arima
from scipy import signal
from scipy.stats import shapiro
from scipy.stats import boxcox
from sklearn.preprocessing import StandardScaler
from scipy.stats import pearsonr
from statsmodels.tsa.api import VAR
from statsmodels.tsa.stattools import grangercausalitytests

import warnings
warnings.filterwarnings("ignore")

[2]: gold = pd.read_excel('gold.xls')
gold
```

```
[2]:
```

	Date	Gold
0	1994-01-01	137.31491
1	1994-02-01	136.52517
2	1994-03-01	133.95854
3	1994-04-01	134.74827
4	1994-05-01	133.46496
..	...	...
347	2022-12-01	617.17670
348	2023-01-01	638.79566
349	2023-02-01	662.68509
350	2023-03-01	634.15597
351	2023-04-01	693.48470

[352 rows x 2 columns]

```
[3]: gold['Date'] = pd.to_datetime(gold['Date'])
gold.set_index('Date', inplace=True)
```

```
[4]: gold.head()
```

```
[4]:
```


	Date	Gold
	1994-01-01	137.31491
	1994-02-01	136.52517
	1994-03-01	133.95854
	1994-04-01	134.74827
	1994-05-01	133.46496

## 1 Time Series Plot


```
[64]: plt.figure(figsize=(8,6))
plt.plot(gold['Gold'])
plt.xlabel('Time')
plt.ylabel('Value')
plt.title('Export Gold Index over Time')
plt.show()
```



```
[6]: plot_acf(gold['Gold']);
```



```
[7]: plot_pacf(gold['Gold']);
```





## 2 STL decomposition

```
[114]: # Multiplicative Decomposition
decomposeM = seasonal_decompose(gold["Gold"],model='multiplicative',
    ↪extrapolate_trend='freq')
plt.rcParams['figure.figsize'] = (12, 8);
#decomposeM.plot();
decomposeM.plot().suptitle('Multiplicative Decomposition', fontsize=16)

# Additive Decomposition
decomposeA = seasonal_decompose(gold["Gold"],model='additive',
    ↪extrapolate_trend='freq')
plt.rcParams['figure.figsize'] = (12, 8);
#decomposeA.plot();
decomposeA.plot().suptitle('Additive Decomposition', fontsize=16)
```

```
[114]: Text(0.5, 0.98, 'Additive Decomposition')
```





```
[9]: ljung_p = np.mean(ljung(x=decomposeA.resid.dropna())['lb_pvalue'])
ljung_p = round(ljung_p, 3)
print("Ljung Box (A), p value:", ljung_p, ", Residuals are uncorrelated" if
      ↪ljung_p>0.05 else ", Residuals are correlated")
```

Ljung Box (A), p value: 0.0 , Residuals are correlated

```
[10]: ljung_p = np.mean(ljung(x=decomposeM.resid.dropna())['lb_pvalue'])
ljung_p = round(ljung_p, 3)
print("Ljung Box (A), p value:", ljung_p, ", Residuals are uncorrelated" if
      ↪ljung_p>0.05 else ", Residuals are correlated")
```

Ljung Box (A), p value: 0.0 , Residuals are correlated

```
[1]: import pandas as pd
import numpy as np
from datetime import datetime
import matplotlib.pyplot as plt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.api import VAR
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.stattools import grangercausalitytests
from scipy.stats import pearsonr
from sklearn.metrics import mean_squared_error
```

### 3 VAR Model

```
[2]: df = pd.read_excel("G&O.xls", index_col=0)
df.head()
```

```
[2]:
```

	Oil	Gold
Date		
1994-01-01	55.11754	137.31491
1994-02-01	52.44123	136.52517
1994-03-01	53.45389	133.95854
1994-04-01	61.19349	134.74827
1994-05-01	66.18445	133.46496

```
[3]: plt.figure(figsize=(12,6))
gold, = plt.plot(df['Gold'])
oil, = plt.plot(df['Oil'], color='red')

for year in range(1994, 2025):
    plt.axvline(datetime(year,1,1), linestyle='--', color='k', alpha=0.1)

plt.xlabel('Time')
plt.ylabel('Price Index')
plt.legend(['Gold', 'Oil'], fontsize=16)
plt.show()
```



### 3.1 ADF Test

```
[4]: adf_test_oil = adfuller(df['Oil'])
adf_statistic_oil = adf_test_oil[0]
adf_pvalue_oil = adf_test_oil[1]
adf_critical_values_oil = adf_test_oil[4]

print('ADF Test for Oil:')
print(f'ADF Statistic: {adf_statistic_oil}')
print(f'p-value: {adf_pvalue_oil}')
```

ADF Test for Oil:  
ADF Statistic: -2.3964936369495327  
p-value: 0.1427285004217443

```
[5]: adf_test_gold = adfuller(df['Gold'])
adf_statistic_gold = adf_test_gold[0]
adf_pvalue_gold = adf_test_gold[1]
adf_critical_values_gold = adf_test_gold[4]

print('ADF Test for Gold:')
print(f'ADF Statistic: {adf_statistic_gold}')
print(f'p-value: {adf_pvalue_gold}')
```

ADF Test for Gold:  
ADF Statistic: 0.3405591027490211  
p-value: 0.9791306790222112



### 3.2 Normalized

```
[6]: # Normalized
avg = df.mean()
dev = df.std()

for col in df:
    df[col] = (df[col]-avg.loc[col]) / dev.loc[col]

plt.figure(figsize=(12,6))
gold, = plt.plot(df['Gold'])
oil, = plt.plot(df['Oil'], color='red')

for year in range(1994, 2025):
    plt.axvline(datetime(year,1,1), linestyle='--', color='k', alpha=0.1)

plt.xlabel('Time')
plt.legend(['Gold', 'Oil'], fontsize=16)
plt.show()
```



```
[7]: # Remove Trend
df = df.diff().dropna()

plt.figure(figsize=(12,6))
gold, = plt.plot(df['Gold'])
oil, = plt.plot(df['Oil'], color='red')
```

```

for year in range(1994, 2025):
    plt.axvline(datetime(year,1,1), linestyle='--', color='k', alpha=0.1)

plt.xlabel('Time')
plt.ylabel('First Difference')
plt.legend(['Gold', 'Oil'], fontsize=16)
plt.show()

```



```

[8]: # Remove Increasing Volatility
vol = df.groupby(df.index.year).std()

df['gold_vol'] = df.index.map(lambda d: vol.loc[d.year, 'Gold'])
df['oil_vol'] = df.index.map(lambda d: vol.loc[d.year, 'Oil'])


df['Gold'] = df['Gold']/df['gold_vol']
df['Oil'] = df['Oil']/df['oil_vol']

plt.figure(figsize=(12,6))
gold, = plt.plot(df['Gold'])
oil, = plt.plot(df['Oil'], color='red')

for year in range(1994, 2025):
    plt.axvline(datetime(year,1,1), linestyle='--', color='k', alpha=0.1)


plt.xlabel('Time')
plt.ylabel('First Difference')
plt.legend(['Gold', 'Oil'], fontsize=16)
plt.show()

```



```
[9]: fig, ax = plt.subplots(figsize=(15, 10))
plot_pacf(df["Gold"], method='ywm', ax=ax)
plt.show
```

```
[9]: <function matplotlib.pyplot.show(close=None, block=None)>
```



### 3.3 VAR Model

```
[10]: var_df = df[["Gold","Oil"]]
var_df.index = pd.DatetimeIndex(var_df.index, freq='MS')

train_data = var_df[:-12]
test_data = var_df[-12:]
```

```
[11]: var_mod = VAR(train_data)
var_mod_fit = var_mod.fit(maxlags=12)
var_mod_fit.summary()
```

[11]: Summary of Regression Results

```
=====
Model:                                VAR
Method:                               OLS
Date:          Thu, 08, Jun, 2023
Time:                                18:30:28
```

```
-----
No. of Equations:      2.00000    BIC:                0.767051
Nobs:                  327.000    HQIC:               0.418778
Log likelihood:        -908.650    FPE:              1.20701
AIC:                   0.187547    Det(Omega_mle):    1.04165
-----
```

Results for equation Gold

```
=====

```

	coefficient	std. error	t-stat	prob
const	0.063003	0.059742	1.055	0.292
L1.Gold	0.135864	0.057532	2.362	0.018
L1.Oil	0.021502	0.058715	0.366	0.714
L2.Gold	-0.053441	0.057868	-0.924	0.356
L2.Oil	0.140323	0.058286	2.407	0.016
L3.Gold	0.030404	0.057641	0.527	0.598
L3.Oil	-0.036469	0.059061	-0.617	0.537
L4.Gold	-0.007928	0.057692	-0.137	0.891
L4.Oil	-0.071068	0.059262	-1.199	0.230
L5.Gold	0.057933	0.057798	1.002	0.316
L5.Oil	-0.040417	0.059400	-0.680	0.496
L6.Gold	-0.009281	0.057661	-0.161	0.872
L6.Oil	-0.014732	0.059161	-0.249	0.803
L7.Gold	0.103544	0.057954	1.787	0.074
L7.Oil	0.001673	0.059025	0.028	0.977
L8.Gold	-0.022509	0.058064	-0.388	0.698
L8.Oil	-0.047345	0.059137	-0.801	0.423
L9.Gold	0.044247	0.058091	0.762	0.446

```
-----
```

L9.Oil	-0.013884	0.058929	-0.236	0.814
L10.Gold	0.123153	0.058298	2.112	0.035
L10.Oil	-0.014806	0.058412	-0.253	0.800
L11.Gold	0.071997	0.057829	1.245	0.213
L11.Oil	0.066070	0.058104	1.137	0.255
L12.Gold	-0.026091	0.057605	-0.453	0.651
L12.Oil	0.006586	0.058581	0.112	0.910

Results for equation Oil

	coefficient	std. error	t-stat	prob
const	0.061039	0.058229	1.048	0.295
L1.Gold	0.029253	0.056075	0.522	0.602
L1.Oil	0.045530	0.057228	0.796	0.426
L2.Gold	0.094658	0.056403	1.678	0.093
L2.Oil	-0.075694	0.056810	-1.332	0.183
L3.Gold	0.108027	0.056181	1.923	0.055
L3.Oil	0.007438	0.057565	0.129	0.897
L4.Gold	0.068013	0.056231	1.210	0.226
L4.Oil	-0.068223	0.057761	-1.181	0.238
L5.Gold	0.017901	0.056334	0.318	0.751
L5.Oil	-0.036593	0.057896	-0.632	0.527
L6.Gold	0.082861	0.056201	1.474	0.140
L6.Oil	0.039106	0.057663	0.678	0.498
L7.Gold	-0.014894	0.056486	-0.264	0.792
L7.Oil	-0.087961	0.057530	-1.529	0.126
L8.Gold	-0.069778	0.056594	-1.233	0.218
L8.Oil	-0.021524	0.057640	-0.373	0.709
L9.Gold	0.020237	0.056620	0.357	0.721
L9.Oil	0.021802	0.057436	0.380	0.704
L10.Gold	0.035720	0.056822	0.629	0.530
L10.Oil	0.043842	0.056933	0.770	0.441
L11.Gold	-0.041313	0.056365	-0.733	0.464
L11.Oil	0.122108	0.056632	2.156	0.031
L12.Gold	-0.063838	0.056146	-1.137	0.256
L12.Oil	-0.077052	0.057098	-1.349	0.177

Correlation matrix of residuals

	Gold	Oil
Gold	1.000000	0.000252
Oil	0.000252	1.000000

Comment:

Therefore our final model will be:

$$Gold = 0.136Gold_{t-1} + 0.14Oil_{t-2}$$

### 3.4 Granger-Causality Test

```
[12]: # Perform Granger-Causality test
max_lag = 5 # Maximum lag order
test_results = grangercausalitytests(var_df, maxlag=max_lag)

# Interpret the results
for lag in range(1, max_lag+1):
    print(f'Granger-Causality test (lag={lag}):')
    result = test_results[lag][0]
    p_value = result['ssr_chi2test'][1]
    if p_value < 0.05:
        print(f' There is evidence of Granger causality (p-value={p_value:.4f})')
    else:
        print(f' No significant evidence of Granger causality_
        (p-value={p_value:.4f})')
    print()
```

Granger Causality

number of lags (no zero) 1

ssr based F test: F=1.1795 , p=0.2782 , df\_denom=347, df\_num=1

ssr based chi2 test: chi2=1.1897 , p=0.2754 , df=1

likelihood ratio test: chi2=1.1877 , p=0.2758 , df=1

parameter F test: F=1.1795 , p=0.2782 , df\_denom=347, df\_num=1

Granger Causality

number of lags (no zero) 2

ssr based F test: F=3.3004 , p=0.0380 , df\_denom=344, df\_num=2

ssr based chi2 test: chi2=6.6968 , p=0.0351 , df=2

likelihood ratio test: chi2=6.6333 , p=0.0363 , df=2

parameter F test: F=3.3004 , p=0.0380 , df\_denom=344, df\_num=2

Granger Causality

number of lags (no zero) 3

ssr based F test: F=2.2737 , p=0.0798 , df\_denom=341, df\_num=3

ssr based chi2 test: chi2=6.9611 , p=0.0731 , df=3

likelihood ratio test: chi2=6.8923 , p=0.0754 , df=3

parameter F test: F=2.2737 , p=0.0798 , df\_denom=341, df\_num=3

Granger Causality

number of lags (no zero) 4

ssr based F test: F=2.0311 , p=0.0897 , df\_denom=338, df\_num=4

```

ssr based chi2 test:  chi2=8.3407 , p=0.0799 , df=4
likelihood ratio test: chi2=8.2420 , p=0.0831 , df=4
parameter F test:      F=2.0311 , p=0.0897 , df_denom=338, df_num=4

```

#### Granger Causality

number of lags (no zero) 5

```

ssr based F test:      F=1.5063 , p=0.1872 , df_denom=335, df_num=5
ssr based chi2 test:  chi2=7.7789 , p=0.1688 , df=5
likelihood ratio test: chi2=7.6928 , p=0.1740 , df=5
parameter F test:      F=1.5063 , p=0.1872 , df_denom=335, df_num=5

```

Granger-Causality test (lag=1):

No significant evidence of Granger causality (p-value=0.2754)

Granger-Causality test (lag=2):

There is evidence of Granger causality (p-value=0.0351)

Granger-Causality test (lag=3):

No significant evidence of Granger causality (p-value=0.0731)

Granger-Causality test (lag=4):

No significant evidence of Granger causality (p-value=0.0799)

Granger-Causality test (lag=5):

No significant evidence of Granger causality (p-value=0.1688)

```
[13]: gc_mod_fit = var_mod.fit(maxlags=12,ic="aic")
```

```
[14]: gc_mod_fit.test_causality("Gold","Oil",kind='f').summary()
```

```
[14]: <class 'statsmodels.iolib.table.SimpleTable'>
```

```
[15]: gc_mod_fit.test_causality("Oil","Gold",kind='f').summary()
```

```
[15]: <class 'statsmodels.iolib.table.SimpleTable'>
```

### 3.5 Forecast

```
[16]: org_df = pd.read_excel("G&O.xls",index_col=0)
      org_df.tail()
```

```
[16]:
```

	Oil	Gold
Date		
2022-12-01	289.90958	617.17670
2023-01-01	285.53345	638.79566
2023-02-01	278.04702	662.68509
2023-03-01	273.70705	634.15597
2023-04-01	277.68535	693.48470

```
[17]: # Forecast 12 steps ahead
forecast = var_mod_fit.forecast(var_mod_fit.endog, steps=12)

forecast_gold = forecast[:, 0]

# Print forecasted values
print("Forecasted Gold Price:", forecast_gold)

Forecasted Gold Price: [-0.23615863 -0.10639113 -0.06046261  0.1164994
0.06634034  0.0424161
0.04570182  0.16772837  0.39596913  0.3107158  0.05819648 -0.00217214]

[18]: # Transform Back
forecast_gold = forecast_gold * np.array(df['gold_vol'][-12:])
forecast_gold = pd.DataFrame(forecast_gold).cumsum()
forecast_gold = forecast_gold*dev[1]+2.1*avg[1]
forecast_gold = np.array(forecast_gold).flatten()
forecast_gold

[18]: array([654.71496905, 652.1338232 , 650.666945 , 653.49332665,
655.10280372, 656.13185712, 657.24062491, 661.3098683 ,
675.61629665, 686.84250823, 688.94515646, 688.8666768 ])


[19]: # Calculate confidence intervals for the forecasted values
alpha = 0.5 # significance level
forecast_ci = var_mod_fit.forecast_interval(var_mod_fit.endog[-var_mod_fit.k_ar:
↵], steps=12, alpha=alpha)

# Access confidence intervals for Gold and Transform Back
ci_lower_gold = pd.DataFrame(forecast_ci[1][:, 0]*np.array(df['gold_vol'][-12:
↵])).cumsum()*dev[1]+2.1*avg[1]
ci_lower_gold = np.array(ci_lower_gold).flatten()
ci_upper_gold = pd.DataFrame(forecast_ci[2][:, 0]*np.array(df['gold_vol'][-12:
↵])).cumsum()*dev[1]+2.1*avg[1]
ci_upper_gold = np.array(ci_upper_gold).flatten()

[29]: # Plot forecast with error bands
plt.figure(figsize=(14, 5))
plt.plot(org_df['Gold'], label='Actual Gold Price')
forecast_dates = pd.date_range(start=org_df.index[-12] + pd.DateOffset(days=1),
↵periods=12, freq='M')
plt.plot(forecast_dates, forecast_gold, label='VAR Forecast Gold Price')
plt.fill_between(forecast_dates,
ci_lower_gold,
ci_upper_gold,
alpha=0.05, label='Error Bands')
plt.legend(fontsize=16, loc='upper left')
```



```
plt.xlabel('Date')
plt.ylabel('Price')
plt.title('VAR Data and Forecast with Error Bands--Gold')
plt.show()
```



### 3.6 Model Performance

```
[24]: # Calculate RMSE
rmse = np.sqrt(mean_squared_error(org_df["Gold"][-12:], forecast_gold))
print("RMSE:", rmse)
```

RMSE: 47.183848206899654

```
[25]: # Calculate MAPE
mape = np.mean(np.abs((org_df["Gold"][-12:] - forecast_gold) /
    ↪ org_df["Gold"][-12:])) * 100
print("MAPE:", mape)
```

MAPE: 6.62067276982332

```
In [1]: ##### Working with time series data in R
library(forecast)
library(ggplot2)
library(readxl)
data = read_excel('G&O.xls')
```

```
Registered S3 method overwritten by 'quantmod':
  method      from
as.zoo.data.frame zoo
```

```
In [2]: library(readxl)
```

```
In [3]: head(data)

data[1,1]
```

A tibble: 6 × 3

Date	Oil	Gold
<dtm>	<dbl>	<dbl>
2008-02-01	369.5281	307.5025
2008-03-01	402.4522	318.6574
2008-04-01	429.2582	330.2073
2008-05-01	483.2078	312.3396
2008-06-01	517.9212	305.5281
2008-07-01	521.1387	307.0089

A tibble: 1 ×

1

Date
<dtm>
2008-02-01

```
In [4]: # convert data to a monthly time series, starting 1/2004
my_ts <- ts(data$Gold, start=1994, freq=12)
my_ts
```

A Time Series: 16 × 12

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep
<b>1994</b>	307.5025	318.6574	330.2073	312.3396	305.5281	307.0089	323.0997	288.2527	287.5617
<b>1995</b>	297.2359	325.1728	318.3613	307.1076	320.3356	325.8638	322.2113	327.4432	343.6328
<b>1996</b>	385.0938	377.4926	383.5143	395.4590	415.1037	424.5804	410.4640	419.0523	438.6968
<b>1997</b>	467.4235	473.2478	490.6219	509.5755	520.5331	525.3702	544.1264	606.5153	605.2320
<b>1998</b>	573.0503	599.4077	576.2093	568.4107	547.2853	550.5429	548.8648	561.9941	602.3692
<b>1999</b>	575.3208	560.6120	548.5686	513.5242	488.1540	461.5992	444.9161	465.9427	463.7710
<b>2000</b>	429.8124	449.2596	460.1185	447.6802	430.3060	455.6762	448.0750	437.1175	418.0652
<b>2001</b>	435.1431	414.3139	416.8806	411.1550	409.0819	403.6525	376.7029	390.4245	395.1628
<b>2002</b>	393.5834	430.5035	425.1728	444.4225	421.3228	468.1145	467.7196	462.3889	440.6713
<b>2003</b>	417.1767	428.9240	432.5765	429.8124	441.8559	421.9151	437.1175	462.4877	438.8942
<b>2004</b>	457.0582	457.3544	460.1185	452.1224	447.6802	432.4778	420.8292	413.8203	414.6107
<b>2005</b>	454.0967	446.0020	446.1007	442.5469	459.8223	483.1194	512.9319	529.3188	516.5842
<b>2006</b>	542.2507	567.8184	556.8608	588.8450	592.3001	613.8203	697.1372	675.2221	658.5390
<b>2007</b>	631.3919	594.1757	603.3564	625.5676	655.0839	626.1599	625.3702	627.5420	609.9702
<b>2008</b>	624.5805	667.4235	669.7927	649.5558	641.6584	608.0948	613.0306	591.6091	587.8578
<b>2009</b>	662.6851	634.1560	693.4847						

In [5]:

```
plot(my_ts,col = 'red', main = 'Gold', ylab = 'Index')


autoplot(my_ts) +
  ggtitle('Gold Price Trends') +
  ylab('Index') +
  xlab('Date')

# stats
mean(my_ts)

sd(my_ts)

median(my_ts)

summary(my_ts)
```




479.892004535519

104.875175410121

454.39289

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
261.6	414.9	454.4	479.9	573.7	697.1

Gold Price Trends



```
In [6]: # assignment
tsmean = mean(my_ts)
print(tsmean)

tsmean <- mean(my_ts)
print(tsmean)

mean(my_ts) -> tsmean
print(tsmean)
```

```
[1] 479.892
[1] 479.892
[1] 479.892
```

## 4. MAPA

```
In [7]: install.packages("MAPA")
```

The downloaded binary packages are in  
/var/folders/81/jjv33vws4cs7b2lsyrkgf1q00000gn/T//Rtmp8wh70B/  
downloaded\_packages

```
In [8]: install.packages("thief")
```

The downloaded binary packages are in  
/var/folders/81/jjv33vws4cs7b2lsyrkgf1q00000gn/T//Rtmp8wh70B/  
downloaded\_packages

```
In [9]: ##### ARIMA #####
library(forecast)
# arima example
model1 = auto.arima(my_ts)
model1

# make a 12 step forecast
fcast = forecast(model1, 12)
fcast

# get the forecasted values
fcast$mean
```

Series: my\_ts  
ARIMA(0,1,0) with drift

Coefficients:  
drift  
2.1208  
s.e. 1.4659

sigma^2 = 393.2: log likelihood = -801.42  
AIC=1606.83 AICc=1606.9 BIC=1613.24

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2009	695.6055	670.1921	721.0189	656.7390	734.4720
May 2009	697.7263	661.7862	733.6663	642.7608	752.6918
Jun 2009	699.8470	655.8297	743.8644	632.5283	767.1658
Jul 2009	701.9678	651.1410	752.7947	624.2349	779.7008
Aug 2009	704.0886	647.2625	760.9148	617.1805	790.9967
Sep 2009	706.2094	643.9595	768.4593	611.0063	801.4124
Oct 2009	708.3302	641.0926	775.5678	605.4991	811.1612
Nov 2009	710.4510	638.5709	782.3310	600.5199	820.3820
Dec 2009	712.5717	636.3314	788.8120	595.9723	829.1712
Jan 2010	714.6925	634.3282	795.0568	591.7859	837.5991
Feb 2010	716.8133	632.5265	801.1001	587.9078	845.7188
Mar 2010	718.9341	630.8994	806.9688	584.2966	853.5715

A Time Series: 2 × 12

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep
<b>2009</b>				695.6055	697.7263	699.8470	701.9678	704.0886	706.2094
<b>2010</b>	714.6925	716.8133	718.9341						

In [10]:

```
##### MAPA example #####
library(MAPA)

# Detailed view of the data at each temporal aggregation level (daily,
# Note: paral = 2 --> Run in parallel cluster mode with e.g., in my ca
mapasimple(my_ts,outplot=2,paral=2)

#Dynamic Fit to the data
mapafit <- mapaest(my_ts,outplot=2, paral=2)
plot(mapafit)
mapafit
mapafor(my_ts,mapafit,ifh=12,fh=0)

# Best fit MAPA model by temporal decomposition + Forecast
# N=None, A=Additive, M=Multiplicative, d=Damped
# Exmaple: MAM = Holt-Winters
mapafit <- mapaest(my_ts,paral=2)
mapafor(my_ts,mapafit)

# Forecast with Error Bands
mapa(my_ts,conf.lvl=c(0.8,0.9,0.95,0.99),outplot=1)
```

Loading required package: parallel

Loading required package: RColorBrewer

Loading required package: smooth

Loading required package: greybox

Package "greybox", v1.0.8 loaded.

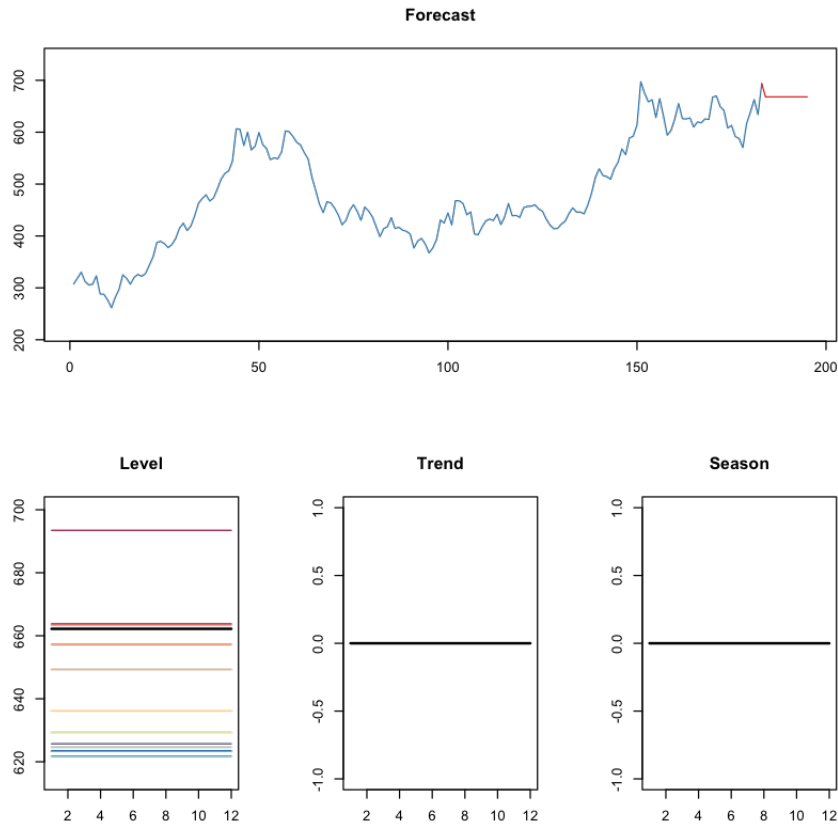
This is package "smooth", v3.2.1

Running with 8 cores

	t+1	t+2	t+3	t+4	t+5	t+6	t+7	t
+8								
668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484
84								
	t+9	t+10	t+11	t+12				
668.0484	668.0484	668.0484	668.0484	668.0484				

Running with 8 cores





MAPA fitted using ets      Original frequency: 12

Aggregation level: 1	Method: ETS(MNN)
Aggregation level: 2	Method: ETS(MNN)
Aggregation level: 3	Method: ETS(MNN)
Aggregation level: 4	Method: ETS(MNN)
Aggregation level: 5	Method: ETS(MNN)
Aggregation level: 6	Method: ETS(MNN)
Aggregation level: 7	Method: ETS(MNN)
Aggregation level: 8	Method: ETS(ANN)
Aggregation level: 9	Method: ETS(ANN)
Aggregation level: 10	Method: ETS(ANN)
Aggregation level: 11	Method: ETS(ANN)
Aggregation level: 12	Method: ETS(ANN)

MAPA fit MSE: 560.89, MAE: 18.06

MAPA fit MSE: 987.74, MAE: 24.14

MAPA fit MSE: 1445.63, MAE: 30.06

MAPA fit MSE: 1898.27, MAE: 34.8

MAPA fit MSE: 2352.04, MAE: 38.85

MAPA fit MSE: 2843.24, MAE: 43.16

MAPA fit MSE: 3367.11, MAE: 45.98

MAPA fit MSE: 3880.34, MAE: 48.85

MAPA fit MSE: 4387.12, MAE: 51.75

MAPA fit MSE: 4897.57, MAE: 54.7

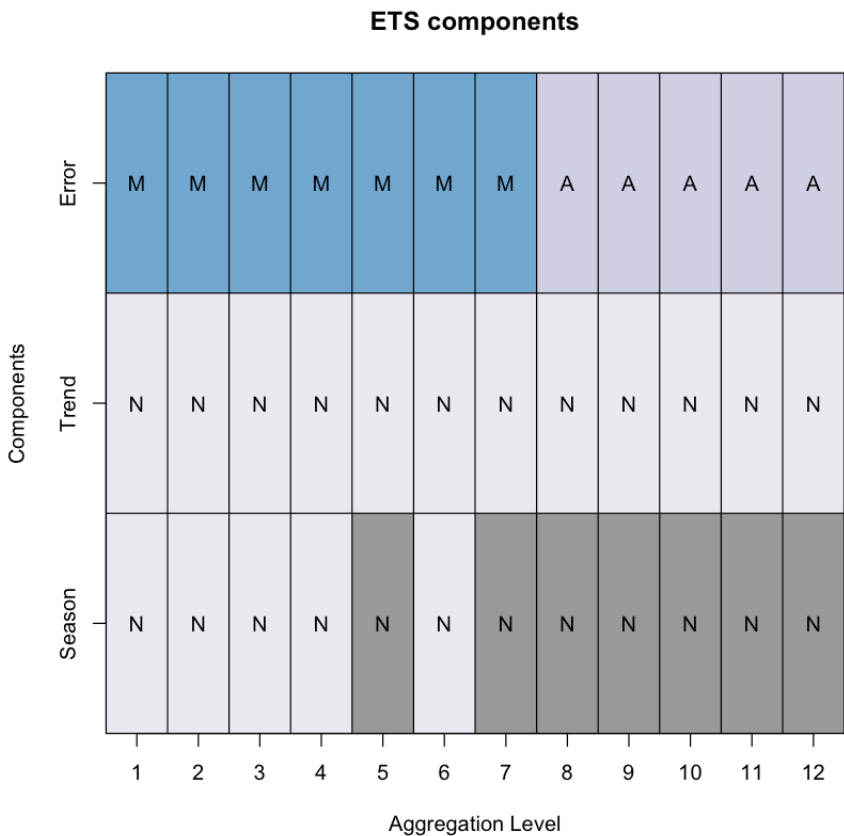
MAPA fit MSE: 5522.04, MAE: 58.64

MAPA fit MSE: 6158.1, MAE: 61.96  
Out-of-sample forecasts:  
NULL

Running with 8 cores

MAPA fit MSE: 560.89, MAE: 18.06  
Out-of-sample forecasts:

	t+1	t+2	t+3	t+4	t+5	t+6	t+7	t+8	t+9	t+10	t+11	t+12
668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484





MAPA fit MSE: 560.89, MAE: 18.06  
Out-of-sample forecasts:

	Lower 0.80	Lower 0.90	Lower 0.95	Lower 0.99	Forecast	Upper 0.99
Upper 0.95						
t+1	637.6973	629.0931	621.6304	607.0447	668.0484	729.0520
	714.4664					
t+2	627.7713	616.3533	606.4499	587.0942	668.0484	749.0026
	729.6469					
t+3	619.3219	605.5086	593.5277	570.1116	668.0484	765.9852
	742.5691					
t+4	610.0400	590.0000	580.0000	555.0000	668.0484	780.0000

t+4	612.2123	596.3835	582.6544	555.8211	668.0484	780.2751
	753.4423					
t+5	605.8960	588.2766	572.9944	543.1263	668.0484	792.9705
	763.1023					
t+6	599.7134	580.3414	563.5391	530.6999	668.0484	805.3969
	772.5577					
t+7	593.6841	572.6029	554.3180	518.5814	668.0484	817.5154
	781.7788					
t+8	588.2175	565.5866	545.9576	507.5939	668.0484	828.5029
	790.1392					
t+9	583.1644	559.1010	538.2296	497.4375	668.0484	838.6593
	797.8672					
t+10	578.3621	552.9372	530.8850	487.7851	668.0484	848.3117
	805.2118					
t+11	572.8158	545.8186	522.4026	476.6374	668.0484	859.4594
	813.6942					
t+12	567.4805	538.9709	514.2431	465.9139	668.0484	870.1829
	821.8537					
	Upper 0.90	Upper 0.80				
t+1	707.0036	698.3995				
t+2	719.7435	708.3255				
t+3	730.5881	716.7749				
t+4	739.7133	723.8845				
t+5	747.8202	730.2008				
t+6	755.7554	736.3834				
t+7	763.4939	742.4127				
t+8	770.5102	747.8793				
t+9	776.9958	752.9324				
t+10	783.1596	757.7347				
t+11	790.2782	763.2810				
t+12	797.1259	768.6163				

Forecast





```
In [11]: mapafit <- mapaest(my_ts, paral=2)
mapafor(my_ts, mapafit)
```

Running with 8 cores

MAPA fit MSE: 560.89, MAE: 18.06

Out-of-sample forecasts:

	t+1	t+2	t+3	t+4	t+5	t+6	t+7	t+8	t+9	t+10	t+11	t+12
668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484	668.0484

```
In [12]: ts_data <- ts(data$Gold, start=2008, freq=12)
train_data <- window(ts_data, end = c(2022, 3))
test_data <- window(ts_data, start = c(2022, 4))
```

## 5. Thief

```
In [13]: ##### Thief example #####
# install.packages('thief', dependencies = TRUE)
library(thief)
library(ggplot2)


fc <- thief(train_data)
autoplot(fc)

fc <- thief(train_data, usemodel = 'arima')
fc
autoplot(fc)
```


	Jan	Feb	Mar	Apr	May	Jun	Jul
Aug 2022				654.3989	656.5300	659.0074	659.4236
2023	663.1512	665.9324	668.0635	663.6346	665.7657	668.5731	667.6691
2024	671.5462	674.3548	676.4860				


	Sep	Oct	Nov	Dec
2022	661.9726	660.9455	663.0766	663.9502
2023	670.7824	669.3399	671.4710	672.4550
2024				

Forecasts from THieF-ETS



Forecasts from THieF-ARIMA





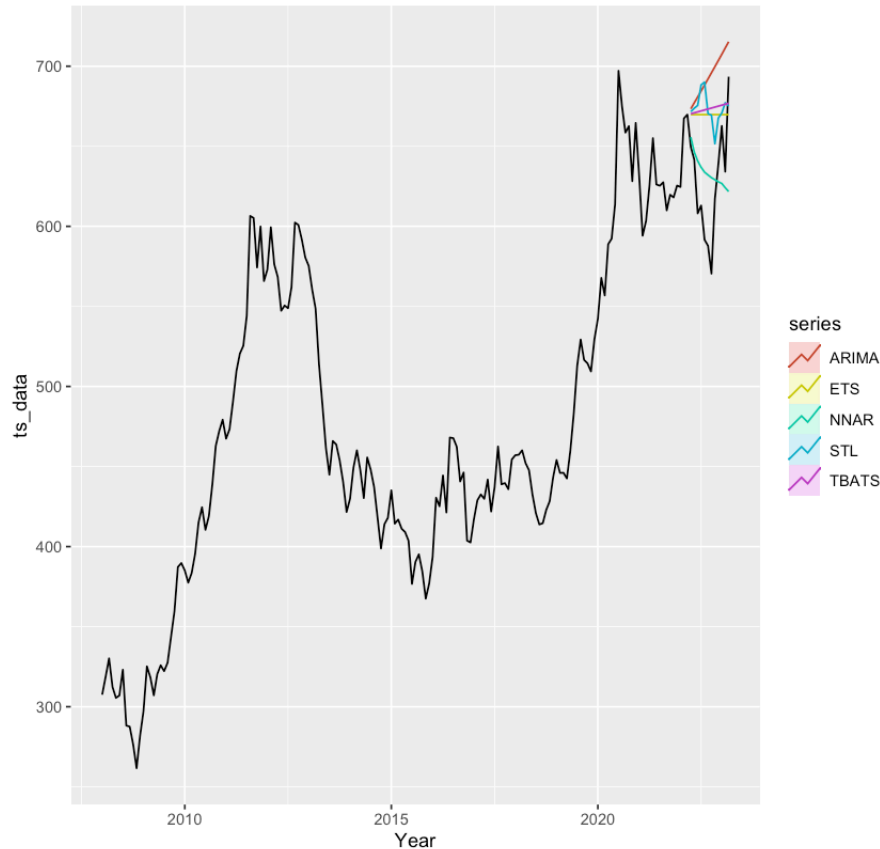
In [14]: `print(mapafit)`

```
MAPA fitted using ets      Original frequency: 12
Aggregation level: 1      Method: ETS(MNN)
Aggregation level: 2      Method: ETS(MNN)
Aggregation level: 3      Method: ETS(MNN)
Aggregation level: 4      Method: ETS(MNN)
Aggregation level: 5      Method: ETS(MNN)
Aggregation level: 6      Method: ETS(MNN)
Aggregation level: 7      Method: ETS(MNN)
Aggregation level: 8      Method: ETS(ANN)
Aggregation level: 9      Method: ETS(ANN)
Aggregation level: 10     Method: ETS(ANN)
Aggregation level: 11     Method: ETS(ANN)
Aggregation level: 12     Method: ETS(ANN)
```


## 6. Combination

```
In [15]: h <- nrow(data) - length(train_data)
ETS <- forecast(ets(train_data), h=h)
ARIMA <- forecast(auto.arima(train_data, lambda=0, biasadj=TRUE),
  h=h)
STL <- stlf(train_data, lambda=0, h=h, biasadj=TRUE)
NNAR <- forecast(nnetar(train_data), h=h)
TBATS <- forecast(tbats(train_data, biasadj=TRUE), h=h)
```

```
In [19]: Combination <- (ETS[["mean"]] + ARIMA[["mean"]] +  
  STL[["mean"]] + NNAR[["mean"]] + TBATS[["mean"]])/5  
autoplot(ts_data) +  
  autolayer(ETS, series="ETS", PI=FALSE) +  
  autolayer(ARIMA, series="ARIMA", PI=FALSE) +  
  autolayer(STL, series="STL", PI=FALSE) +  
  autolayer(NNAR, series="NNAR", PI=FALSE) +  
  autolayer(TBATS, series="TBATS", PI=FALSE) +  
  xlab("Year")
```




```
In [20]: autoplot(ts_data) +  
         autolayer(Combination, series="Combination")
```






```
In [21]: autoplot(ts_data) +  
         autolayer(NNAR, series="NNAR", PI=FALSE)
```



```
In [22]: autoplot(ts_data) +
         autolayer(ARIMA, series="ARIMA", PI=FALSE)
```



## 7. ARIMA

```
In [23]: ARIMA <- forecast(auto.arima(train_data, lambda=0, biasadj=TRUE), h=h)
# Forecast using the trained model
forecast <- forecast(ARIMA, h = 12)

# Extract the forecasted values
forecast_values <- forecast$mean

# Calculate the Mean Absolute Error (MAE)
mae <- mean(abs(forecast_values - test_data))

# Print the MAE
print(mae)
```

```
[1] 68.418
```

In [24]:

```
mse <- mean((forecast_values - test_data)^2)
mrse <- sqrt(mse)
print(mrse)
```

```
[1] 75.06919
```


In [25]:

```
mape <- mean(abs((forecast_values - test_data) / test_data)) * 100
# Print the MAPE
print(mape)
```


```
[1] 11.2063
```

In [26]:

```
fitted_values_arima <- fitted(ARIMA)
plot(fitted_values_arima - train_data, type = "p", ylab = "Residuals")
```



```
In [27]: plot(forecast_values - test_data, type = "p", ylab = "Residuals")
```



In [28]: `summary(ARIMA)`

Forecast method: ARIMA(1,1,0)(1,0,0)[12] with drift

Model Information:

Series: train\_data

ARIMA(1,1,0)(1,0,0)[12] with drift

Box Cox transformation: lambda= 0

Coefficients:

	ar1	sar1	drift
	0.0599	0.0016	0.0046
s.e.	0.0766	0.0811	0.0032

sigma^2 = 0.001591: log likelihood = 307.98

AIC=-607.96 AICc=-607.72 BIC=-595.42

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
Training set	-0.4102585	18.97984	14.7419	-0.154984	3.154251	0.2449689
	ACF1					
Training set	-0.02497386					

Forecasts:

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2022		673.4007	639.3327	708.1569	622.2634	727.5824
May 2022		677.1454	627.4772	728.2826	603.2153	757.5747
Jun 2022		680.8178	619.2381	744.6603	589.7342	781.9150
Jul 2022		684.5578	612.9195	759.2616	579.1500	803.5333
Aug 2022		688.3231	607.7748	772.7479	570.3447	823.4611
Sep 2022		692.0733	603.4141	785.4282	562.7449	842.1907
Oct 2022		695.8927	599.7044	797.6045	556.1040	860.1393
Nov 2022		699.7117	596.4523	809.3310	550.1634	877.4253
Dec 2022		703.5676	593.5985	820.7416	544.8138	894.2339
Jan 2023		707.4290	591.0507	831.8660	539.9307	910.6260
Feb 2023		711.3881	588.8362	842.8619	535.5090	926.7959
Mar 2023		715.2971	586.7941	853.5931	531.3790	942.6104

In [29]:

```

model_a <- forecast(arima(train_data, order = c(1, 1, 0)), h=h)

# Print the model summary
summary(model_a)

```

Forecast method: ARIMA(1,1,0)

Model Information:

Call:

```
arima(x = train_data, order = c(1, 1, 0))
```

Coefficients:

```

      ar1
      0.0471
s.e.    0.0765

```

sigma^2 estimated as 364.2: log likelihood = -742.51, aic = 1489.03

Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
ACF1						
Training set	2.021343	19.02692	14.71198	0.3600966	3.15007	0.2444716
	0.01170269					

Forecasts:

	Point	Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Apr 2022		669.9042	645.4486	694.3597	632.5026	707.3057
May 2022		669.9094	634.5010	705.3179	615.7569	724.0620
Jun 2022		669.9097	626.1808	713.6386	603.0321	736.7872
Jul 2022		669.9097	619.2066	720.6128	592.3660	747.4534
Aug 2022		669.9097	613.0819	726.7375	582.9990	756.8203
Sep 2022		669.9097	607.5559	732.2635	574.5478	765.2716
Oct 2022		669.9097	602.4812	737.3381	566.7868	773.0326
Nov 2022		669.9097	597.7627	742.0567	559.5704	780.2490
Dec 2022		669.9097	593.3343	746.4850	552.7978	787.0216
Jan 2023		669.9097	589.1484	750.6709	546.3960	793.4234
Feb 2023		669.9097	585.1690	754.6503	540.3100	799.5093
Mar 2023		669.9097	581.3683	758.4510	534.4973	805.3220

```
In [30]: forecast <- forecast(model_a, h = 12)

# Extract the forecasted values
forecast_values_a <- forecast$mean

mse <- mean((forecast_values_a - test_data)^2)
mrse <- sqrt(mse)
print(mrse)

[1] 55.26651
```

```
In [31]: mape <- mean(abs((forecast_values_a - test_data) / test_data)) * 100


# Print the MAPE
print(mape)

[1] 7.928594
```

```
In [36]: fitted_values <- ARIMA$mean
```


```
In [37]: residual <- forecast_values - test_data
```

```
In [38]: plot(residual, type = "l", ylab = "Residuals")
```




```
In [39]: acf(residual)

# Plot the PACF
pacf(residual)
```







```
In [40]: fitted_values_arima <- fitted(ARIMA)
```

```
In [42]: mse <- mean((forecast_values - test_data)^2)
mrse <- sqrt(mse)
print(mrse)
```

```
[1] 75.06919
```

```
In [43]: mape <- mean(abs((forecast_values - test_data) / test_data)) * 100

# Print the MAPE
print(mape)
```

```
[1] 11.2063
```

## 8. NNAR

```
In [44]: library(forecast)

# Train the NNAR model
model <- nnetar(train_data)

# Forecast using the trained model
forecast <- forecast(model, h = 12)

# Extract the forecasted values
forecast_values_NNAR <- forecast$mean

# Calculate the Mean Absolute Error (MAE)
mae <- mean(abs(forecast_values_NNAR - test_data))

# Print the MAE
print(mae)

[1] 29.69858
```

```
In [45]: fitted_values_nnar <- fitted(model)
```


```
In [47]: current_dir <- getwd()
print(current_dir)

[1] "/Users/dwx/Documents/spring2023/ECON412 Big Data/Final Project/thief"
```

```
In [48]: mse <- mean((forecast_values_NNAR - test_data)^2)
mrse <- sqrt(mse)
print(mrse)

[1] 36.40232
```


```
In [49]: autoplot(ts_data) +  
         autolayer(NNAR, series="NNAR", PI=FALSE)
```



```
In [50]: mape <- mean(abs((forecast_values_NNAR - test_data) / test_data)) * 100  
  
# Print the MAPE  
print(mape)
```

```
[1] 4.77901
```


```
In [51]: residuals <- train_data-fitted_values_nnar  
plot(residuals, type = "p", ylab = "Residuals")
```



In [52]:

```
residuals <- test_data - forecast_values_NNAR

# Create a residual plot
plot(residuals, type = "p", ylab = "Residuals")
```



## 9. ETS

In [53]:

```
model <- ets(train_data)
forecast <- forecast(model, h=12)
forecast_values_ets <- forecast$mean
forecast_values_ets
```

A Time Series: 2 × 12

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep
<b>2022</b>				669.7925	669.7925	669.7925	669.7925	669.7925	669.7925
<b>2023</b>	669.7925	669.7925	669.7925						

```
In [54]: fitted_values_ets <- fitted(model)
```

```
In [55]: print(model)
```

ETS(M,N,N)

Call:

```
ets(y = train_data)
```

Smoothing parameters:

alpha = 0.9999

Initial states:

l = 307.0221

sigma: 0.0404

	AIC	AICc	BIC
	1880.220	1880.363	1889.645

```
In [56]: aic_score <- AIC(model)
bic_score <- BIC(model)

# Print AIC and BIC scores
cat("AIC:", aic_score, "\n")
cat("BIC:", bic_score, "\n")
```

AIC: 1880.22

BIC: 1889.645

```
In [57]: mse <- mean((forecast_values - test_data)^2)
rmse <- sqrt(mse)
print(rmse)
```


[1] 75.06919

```
In [58]: mape <- mean(abs((forecast_values - test_data) / test_data)) * 100

# Print the MAPE
print(mape)
```

[1] 11.2063

```
In [66]: autoplot(ts_data) +  
         autolayer(ETS, series="ETS", PI=FALSE)
```



## 10. LSTM

```
In [1]: ▶ import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from arch import arch_model
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: ▶ data = pd.read_excel("G&O.xls")
data.set_index('Date', inplace=True)
```

```
In [3]: ▶ data
```

Out[3]:

	Oil	Gold
Date		
1994-01-01	55.11754	137.31491
1994-02-01	52.44123	136.52517
1994-03-01	53.45389	133.95854
1994-04-01	61.19349	134.74827
1994-05-01	66.18445	133.46496
...	...	...
2022-12-01	289.90958	617.17670
2023-01-01	285.53345	638.79566
2023-02-01	278.04702	662.68509
2023-03-01	273.70705	634.15597
2023-04-01	277.68535	693.48470

352 rows × 2 columns

```
In [4]: ▶ train_data = data["Gold"][:-12]
test_data = data["Gold"][-12:]
```



In [5]: `test_data`

Out[5]:

Date	
2022-05-01	649.55577
2022-06-01	641.65844
2022-07-01	608.09477
2022-08-01	613.03060
2022-09-01	591.60908
2022-10-01	587.85785
2022-11-01	570.48371
2022-12-01	617.17670
2023-01-01	638.79566
2023-02-01	662.68509
2023-03-01	634.15597
2023-04-01	693.48470

Name: Gold, dtype: float64

In [6]: `from sklearn.preprocessing import MinMaxScaler`

```
scaler = MinMaxScaler(feature_range=(0, 1))
train_data_scaled = scaler.fit_transform(train_data.values.reshape(-1, 1))
test_data_scaled = scaler.transform(test_data.values.reshape(-1, 1))
```

In [7]: `from tensorflow.keras.models import Sequential`  
`from tensorflow.keras.layers import LSTM, Dense`

```
model = Sequential()
model.add(LSTM(units=50, activation='relu', input_shape=(1, 1)))
model.add(Dense(units=1))
```

```
In [8]: ▶ model.compile(optimizer='adam', loss='mean_squared_error')  
model.fit(train_data_scaled.reshape(-1, 1, 1), train_data_scaled, epochs
```

```
Epoch 1/40
34/34 [=====] - 2s 2ms/step - loss: 0.1716
Epoch 2/40
34/34 [=====] - 0s 2ms/step - loss: 0.1110
Epoch 3/40
34/34 [=====] - 0s 2ms/step - loss: 0.0679
Epoch 4/40
34/34 [=====] - 0s 2ms/step - loss: 0.0443
Epoch 5/40
34/34 [=====] - 0s 2ms/step - loss: 0.0347
Epoch 6/40
34/34 [=====] - 0s 2ms/step - loss: 0.0284
Epoch 7/40
34/34 [=====] - 0s 3ms/step - loss: 0.0224
Epoch 8/40
34/34 [=====] - 0s 3ms/step - loss: 0.0170
Epoch 9/40
34/34 [=====] - 0s 3ms/step - loss: 0.0120
Epoch 10/40
34/34 [=====] - 0s 3ms/step - loss: 0.0080
Epoch 11/40
34/34 [=====] - 0s 3ms/step - loss: 0.0050
Epoch 12/40
34/34 [=====] - 0s 3ms/step - loss: 0.0029
Epoch 13/40
34/34 [=====] - 0s 3ms/step - loss: 0.0016
Epoch 14/40
34/34 [=====] - 0s 3ms/step - loss: 8.7218e-04
Epoch 15/40
34/34 [=====] - 0s 3ms/step - loss: 5.1604e-04
Epoch 16/40
34/34 [=====] - 0s 3ms/step - loss: 3.5598e-04
Epoch 17/40
34/34 [=====] - 0s 3ms/step - loss: 2.9312e-04
Epoch 18/40
34/34 [=====] - 0s 3ms/step - loss: 2.6440e-04
Epoch 19/40
34/34 [=====] - 0s 2ms/step - loss: 2.5467e-04
Epoch 20/40
34/34 [=====] - 0s 2ms/step - loss: 2.4419e-04
Epoch 21/40
34/34 [=====] - 0s 2ms/step - loss: 2.4051e-04
Epoch 22/40
34/34 [=====] - 0s 2ms/step - loss: 2.3575e-04
Epoch 23/40
34/34 [=====] - 0s 2ms/step - loss: 2.3809e-04
Epoch 24/40
34/34 [=====] - 0s 2ms/step - loss: 2.2725e-04
Epoch 25/40
34/34 [=====] - 0s 2ms/step - loss: 2.2605e-04
Epoch 26/40
34/34 [=====] - 0s 2ms/step - loss: 2.2354e-04
Epoch 27/40
34/34 [=====] - 0s 2ms/step - loss: 2.1807e-04
Epoch 28/40
34/34 [=====] - 0s 2ms/step - loss: 2.1431e-04
Epoch 29/40
```

```

34/34 [=====] - 0s 2ms/step - loss: 2.1536e-04
Epoch 30/40
34/34 [=====] - 0s 3ms/step - loss: 2.1189e-04
Epoch 31/40
34/34 [=====] - 0s 3ms/step - loss: 2.0698e-04
Epoch 32/40
34/34 [=====] - 0s 3ms/step - loss: 2.0512e-04
Epoch 33/40
34/34 [=====] - 0s 2ms/step - loss: 1.9947e-04
Epoch 34/40
34/34 [=====] - 0s 2ms/step - loss: 1.9775e-04
Epoch 35/40
34/34 [=====] - 0s 3ms/step - loss: 1.9441e-04
Epoch 36/40
34/34 [=====] - 0s 3ms/step - loss: 1.9238e-04
Epoch 37/40
34/34 [=====] - 0s 3ms/step - loss: 1.8688e-04
Epoch 38/40
34/34 [=====] - 0s 2ms/step - loss: 1.8172e-04
Epoch 39/40
34/34 [=====] - 0s 3ms/step - loss: 1.8507e-04
Epoch 40/40
34/34 [=====] - 0s 3ms/step - loss: 1.7378e-04

```

Out[8]: <keras.callbacks.History at 0x2059b4e90a0>

```

In [16]: ► predictions_scaled = model.predict(train_data_scaled[-12:].reshape(-1, 1))
          predictions = scaler.inverse_transform(predictions_scaled)

```

```

1/1 [=====] - 0s 194ms/step

```

```

In [15]: ► predictions = np.squeeze(predictions)
          predictions = pd.Series(predictions, index=test_data.index)
          predictions

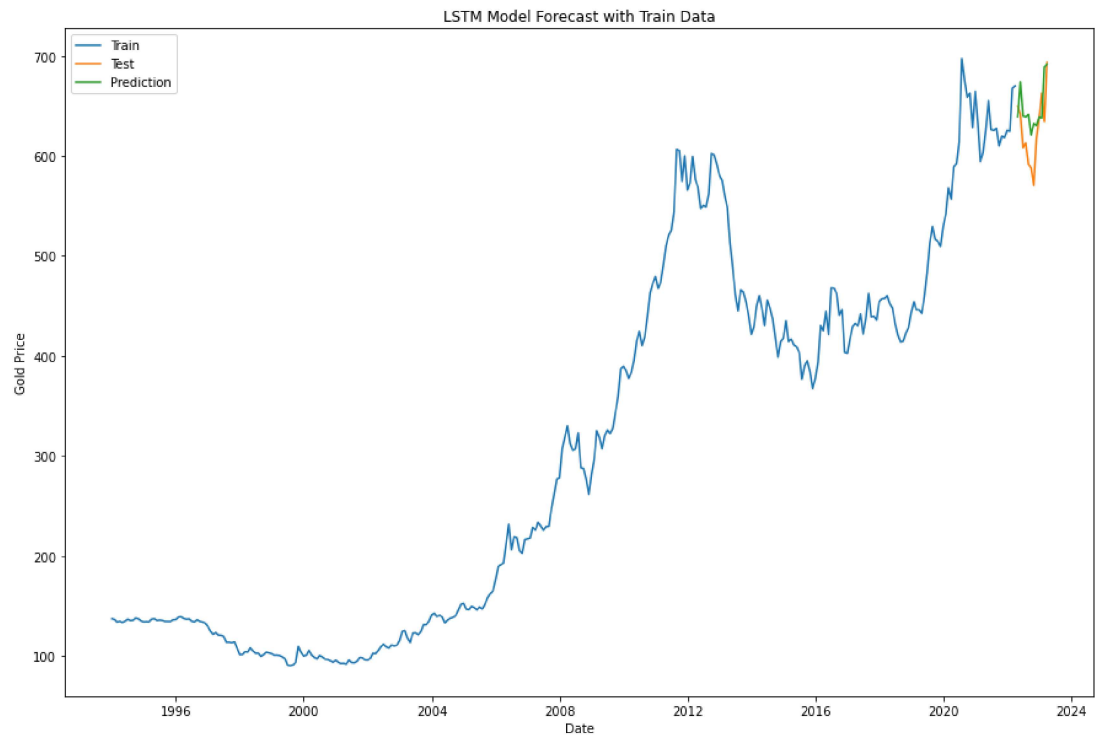
```

Out[15]:


Date	
2022-05-01	638.990967
2022-06-01	673.876282
2022-07-01	639.685547
2022-08-01	638.759399
2022-09-01	641.307068
2022-10-01	620.780640
2022-11-01	632.172791
2022-12-01	630.212219
2023-01-01	638.875122
2023-02-01	637.833801
2023-03-01	688.624817
2023-04-01	691.467590

dtype: float32

```
In [11]: ▶ plt.figure(figsize=(15, 10))
plt.plot(train_data, label='Train')
plt.plot(test_data, label='Test')
plt.plot(predictions, label='Prediction')
plt.legend()
plt.xlabel('Date')
plt.ylabel('Gold Price')
plt.title('LSTM Model Forecast with Train Data')
plt.show()
```



```
In [12]: ▶ plt.figure(figsize=(15, 10))
plt.plot(test_data, label='Test')
plt.plot(predictions, label='Prediction')
plt.legend()
plt.xlabel('Date')
plt.ylabel('Gold Price')
plt.title('LSTM Model Forecast')
plt.show()
```




```
In [13]: ▶ from sklearn.metrics import mean_squared_error
import numpy as np

rmse = np.sqrt(mean_squared_error(test_data, predictions))
mape = np.mean(np.abs((test_data - predictions) / test_data)) * 100
print("RMSE:", rmse)
print("MAPE:", mape)
```

```
RMSE: 34.06774754057272
MAPE: 4.634053137000738
```

```
In [14]: residuals = test_data - predictions
plt.figure(figsize=(10, 6))
plt.scatter(range(len(residuals)), residuals)
plt.axhline(0, color='r', linestyle='--')
plt.xlabel('Index')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()
```



## 11. Reference

1. U.S. Bureau of Labor Statistics, Export Price Index (End Use): Nonmonetary Gold [IQ12260], retrieved from FRED, Federal Reserve Bank of St. Louis; <https://fred.stlouisfed.org/series/IQ12260>, June 9, 2023.
2. U.S. Energy Information Administration, Crude Oil Prices: West Texas Intermediate (WTI) - Cushing, Oklahoma [DCOILWTICO], retrieved from FRED, Federal Reserve Bank of St. Louis; <https://fred.stlouisfed.org/series/DCOILWTICO>, June 9, 2023.

### Coding Source:

- [1] pandas (Version 1.3.5). McKinney, W. Retrieved from <https://pandas.pydata.org>
- [2] Seabold, S. & Perktold, J. Retrieved from <http://conference.scipy.org/proceedings/scipy2010/seabold.html>
- [3] NumPy (Version 1.22.0). Oliphant, T. E. Retrieved from <https://numpy.org>
- [4] Matplotlib (Version 3.5.1). Hunter, J. D.. Retrieved from <https://matplotlib.org>

在此处键入文本